



# Introduction to Deep Reinforcement Learning and Multi-Agent Modelling

Presented

by Thomas Asikis, [asikist@ethz.ch](mailto:asikist@ethz.ch)

**Agent-Based Modeling and Social System Simulation**

**Fall Semester 2019**

# Usual Setting

- Optimization Objective:  $\min(O(a, s)) \vee \max(O(a, s))$ .
- Time  $t$ : Usually discrete, (non-)uniform time intervals.
- State  $s_t \in \mathbb{S} \subset \mathbb{R}^{n \times m \times \dots}$
- Actions  $a_t, \in \mathbb{A} \subset \mathbb{R}^{n \times m \times \dots}$
- Reward:  $r_t = h(O(a, s))$ , usually noise estimate of  $(O(a, s))$
- Environment  $\rightarrow$  State transition:  $s_{t+1} = g(s_t, a_t)$
- Code for slides of the 2<sup>nd</sup> will be uploaded to:  
[https://github.com/asikist-ethz/reinforcement\\_learning](https://github.com/asikist-ethz/reinforcement_learning)

# $\epsilon$ -Greedy Policy

## $\epsilon$ -Greedy Policy Algorithm

- **Inputs:**  $s, Q(s, a)$
- **Parameters:** small  $\epsilon > 0$

#Determine Optimal action

$$a^* \leftarrow \underset{a}{\operatorname{argmax}} Q(s_t, a)$$

**For**  $a$  in  $\mathcal{A}_t$ :  $\# \mathcal{A}_t \leftarrow \text{environment.all\_possible\_actions}(s_t)$

$$\pi(a|s_t) = \begin{cases} 1 - \epsilon + \frac{\epsilon}{|\mathcal{A}_t|}, & a = a^* \\ \frac{\epsilon}{|\mathcal{A}_t|}, & a \neq a^* \end{cases}$$

**Return**  $a \leftarrow \pi(a|s_t)$

# Reward Design

- Rewards on irrelevant tasks: e.g. capture the flag vs disabling an enemy player, over-maximizing a single asset instead of average portfolio.
- Sparse Rewards: Some tasks may be difficult to model as dense rewards, so learning may suffer.
- Reward and goal monotonicity: Rewards needs to increase as we approach the goal.
- Imitation vs Optimization!
- Inverse RL: Look at optimal behaviors and infer the reward.
- Some board examples:

# On-policy Reinforcement Learning

Update a policy based on actions taken by that policy. Exploration is included in the policy learning model.

The policy learned is soft greedy but not greedy (always near optimal)

Example:

1. Chose action  $a$  based on  $\pi$  (e.g.  $\epsilon$ -greedy\*)
2. State transition
3. Update  $\pi$  based on outcome of  $a$

nice discussion: <https://datascience.stackexchange.com/questions/26471/is-my-understanding-of-on-policy-and-off-policy-td-algorithms-correct>

\*choose an action greedily or randomly based on a probability, e.g.  $\epsilon = 0.1$  for random choice, and  $1 - \epsilon = 0.9$  for deterministic. When choosing randomly, usually the actions are chosen uniformly. Then the probability of choosing any action becomes:  $\frac{\epsilon}{|A|}$

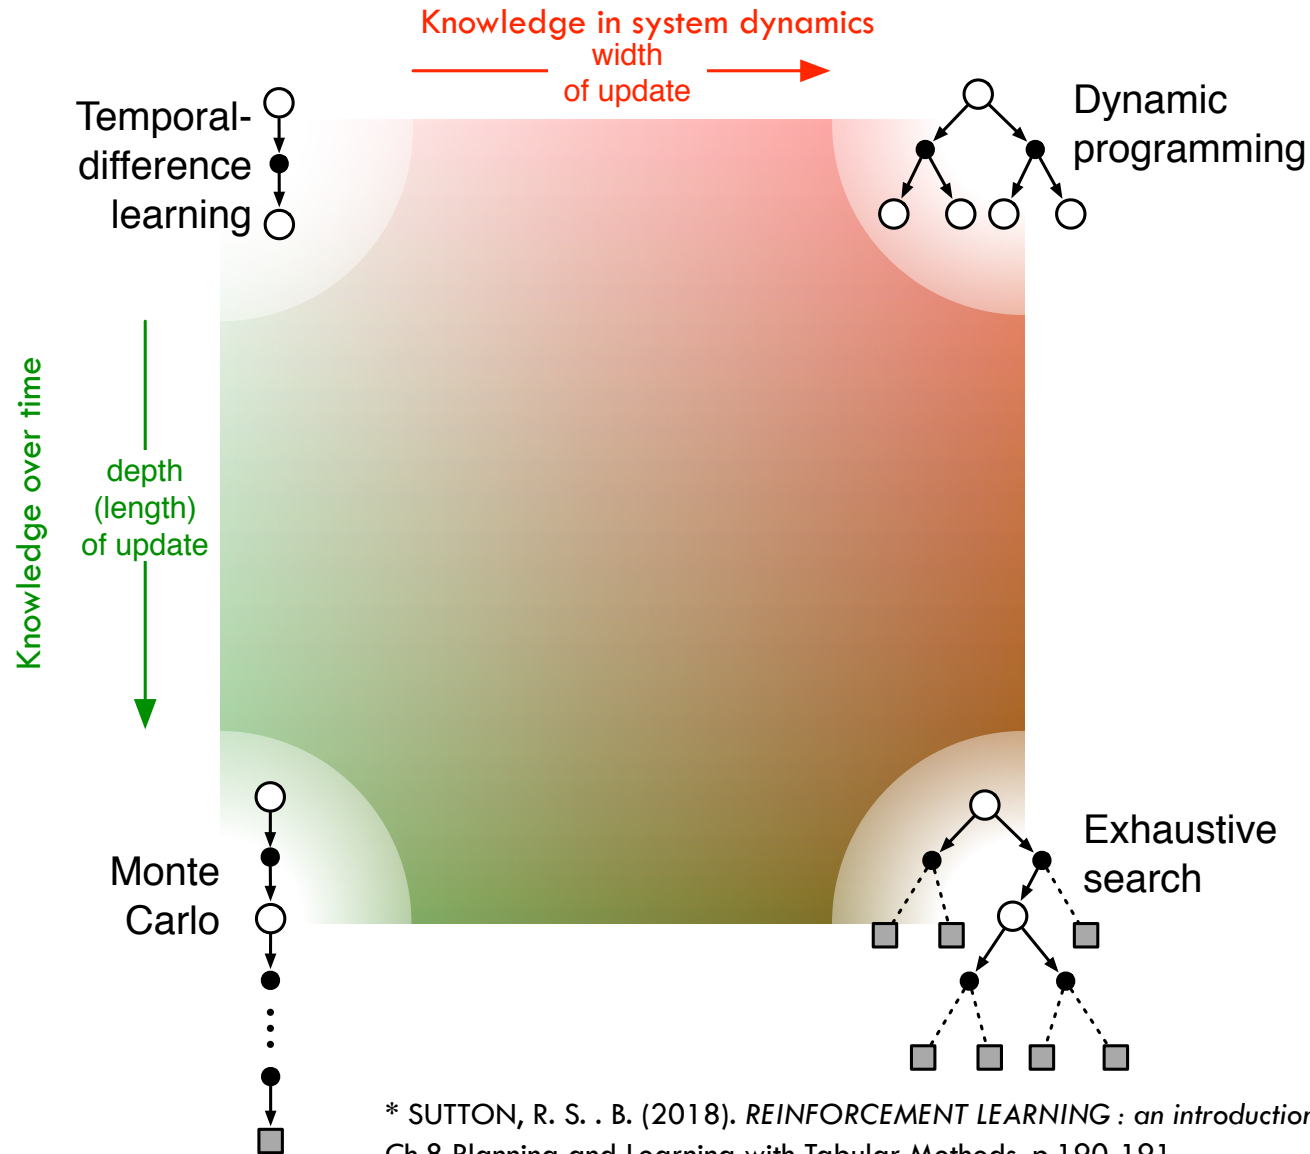
# Off-policy Reinforcement Learning

Update an optimal/greedy policy based on action taken by another exploratory policy

Example:

1. Chose action  $a$  based on  $\pi' \neq \pi$
  2. State transition
  3. Update  $\pi$  based on outcome of  $a$
- \*  $\pi'$ : target policy that is evaluated or learned (e.g.  $\max(Q_{t+1})$ )  
 $\pi$ : behavior policy that affects choice of actions (e.g.  $\epsilon$ -greedy)

# Dimensions of the Reinforcement Learning Problem



\* SUTTON, R. S. . B. (2018). *REINFORCEMENT LEARNING : an introduction*. MIT PRESS.  
Ch.8 Planning and Learning with Tabular Methods, p.190-191

# Temporal Difference

- Combination of Dynamic Programming and Monte-Carlo methods
- On-line application: step by step updates.
- Can be extended to  $n$ -step updates (e.g. eligibility traces)
- System dynamics can be unknown.
- A simple state value TD Method:

$$V(s_t) \leftarrow V(s_t) + \alpha[r_t + \gamma V(S_{t+1}) - V(S_t)]$$

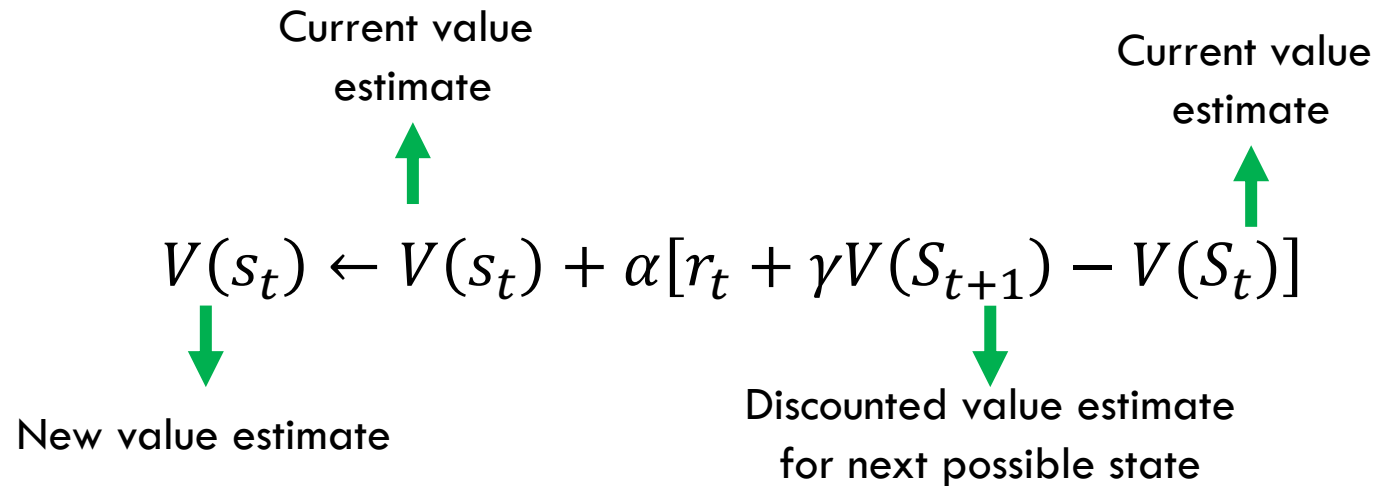
Diagram illustrating the TD update equation with annotations:

- Current value estimate**: Points to  $V(s_t)$  (green arrow pointing up).
- Reward**: Points to  $r_t$  (green arrow pointing up).
- Current value estimate**: Points to  $V(S_t)$  (green arrow pointing up).
- New value estimate**: Points to the new  $V(s_t)$  (green arrow pointing down).
- Learning Rate**: Points to  $\alpha$  (red arrow pointing down).
- Discounted value estimate for next possible state**: Points to  $\gamma V(S_{t+1})$  (green arrow pointing down).



# Bootstrapping

- Updating estimates based on other estimates:



# Q-learning

$$Q(s_t, a_t) \leftarrow (1 - \alpha) Q(s_t, a_t) + \alpha \cdot \left( r_t + \gamma \cdot \max_a [Q(s_{t+1}, a)] \right)$$

Diagram illustrating the Q-learning update equation with annotations:

- Learning rate** (red arrow pointing to  $\alpha$ )
- Discount factor** (red arrow pointing to  $\gamma$ )
- Reward** (green arrow pointing to  $r_t$ )
- Value from previous calculation** (green arrow pointing to  $Q(s_t, a_t)$ )
- $\pi'$ : greedy policy over  $\pi$**  (green arrow pointing to  $\max_a$ )
- Estimate of optimal future value, using  $\epsilon$ -greedy.** (green arrow pointing to  $Q(s_{t+1}, a)$ )

# Q-Learning

## Off Policy TD Control (Q-Learning) Algorithm

- **Inputs:**  $\epsilon$ -greedy  $\pi(a|s, Q(s, a))$ , environment
- **Parameters:** step-size  $\alpha$ , small  $\epsilon > 0$ ,  $0 \leq \gamma \leq 1$
- **Variables:**  $Q(s, a)$ , the expected return after taking and action  $a$  on state  $s$ . Arbitrary initiation.  $Q(s_{terminal}, a) = 0 \forall a$

`total_episodes`

**For** episode **in** `total_episodes`:

$t \leftarrow 0$

$s_t \leftarrow \text{environment.reset}()$  # reset and get initial state

**While**  $t < T$  **and**  $s_t \neq s_{terminal}$ :

$a_t \leftarrow \pi(a|s, Q(s, a))$

$s_{t+1}, r_t \leftarrow \text{environment.step}(a)$

$Q(s_t, a_t) \leftarrow (1 - \alpha) Q(s_t, a_t) + \alpha \cdot (r_t + \gamma \cdot \max_a [Q(s_{t+1}, a)])$

$s_t \leftarrow s_{t+1}$

**Return**  $\pi(a|s, Q(s, a)), Q(s, a)$

## Challenges for classic RL

- Maximum operator can be biased. E.g. stochastic policies, maximization bias etc.
- Decision and state may be non-linearly dependent... More complex estimators than “mean” are required.
- States might be continuous, too many. e.g. if state is a permutation of values... Table methods are inefficient or don't work.
- Decision may need to be done in continuous time – Continuous MDP.
- Partially Observable MDP?

Plugging a good value and policy estimator can tackle most of the above “efficiently”.

# Partially Observable Markov Decision Process

- MDPs are rare in real world scenarios
- State aliasing is often. Same state-actions → completely different rewards.
- Terminal states difficult to identify.
- State transformation or learning can help in this cases. Such Transformation involve a Representation Learning Task.


# On-policy Distribution

$\mu(s)$ :

- The number of timesteps spent in the state  $s$  under a policy  $\pi$ .
- Also called on-policy distribution.
- Can also be parametrized or estimated as well

Intuition: Finding the optimal policy and value may strongly depend from this distribution, which is usually unknown.

Worst case: it changes when policy changes...



**NOTHING TO  
DO HERE**

**Too many challenges**

Thanks for your attention!

# The Prediction Problem: Precise estimation of values

- Parametrize the value function:  $V(s, \mathbf{w})$
- Given a policy  $\pi$ , update parameters  $\mathbf{w}$  s.t.:

$$\min_{\mathbf{w}} (\overline{VE}(\mathbf{w}))$$

Where  $\overline{VE}$  is a prediction objective (e.g. approximation error metric):

$$\overline{VE}(\mathbf{w}) = \sum_{s \in \mathcal{S}} \mu(s) \cdot [V_{\pi}(s) - V(s, \mathbf{w})]^2$$

Iteratively update the parameters with gradient update:

$$\mathbf{w}' = \mathbf{w} - \eta \nabla_{\mathbf{w}} \overline{VE}(\mathbf{w})$$

Empirical Estimation of  $\mu(s)$  can be used here, as the true value.

$\eta$ : as in machine learning, it is the learning rate. Usually a constant in  $(0,1)$ .



# Policy Gradient

- A policy can be parametrized  $\pi(a|s, \theta)$ .
- And also approximated by maximizing some policy performance metric.
- Approximate optimal policy by learning parameters
$$\pi(\theta) \rightarrow \pi^*$$
- Update  $\theta' = \theta + \eta \nabla_{\theta} J(\theta)$  , where  $J = v_{\pi, \theta}(s_o)$  a learning performance metric to maximize.
- Recall from previous lecture:

Each policy has its own true state value function:

$$v_{\pi, \theta}(s) = \mathbb{E}_{\pi}(G_t | S_t = s, \theta)$$

# Policy Gradient

- But  $v_{\pi, \theta}(s_o)$  depends on action selection  $\pi$  and state distribution  $\mu(s)$ .  
Changing the policy parameter  $\theta$  affects both, and thus performance as well:  
Assumption:  $\nabla_{\theta} J(\theta)$  depends on  $\nabla_{\theta} \mu(s|\theta)$
- The rate of change of policy distribution  $\nabla_{\theta} \mu(s|\theta)$  is very difficult to estimate.

# Policy Gradient Theorem

Policy gradient theorem: The learning performance gradient is proportional to  $\mu, Q, \nabla \pi$

$$\nabla_{\theta} J(\theta) \propto \sum_{s \in \mathcal{S}} \mu(s) \sum_{a \in \mathcal{A}} Q_{\pi}(s, a) \nabla_{\theta} \pi(a|s, \theta)$$

- And so it doesn't involve a derivative on the state distribution!!!
- So we can calculate the policy gradient.

# Actor Critic

- Combine policy gradient with value estimation
- Use one set of parameters  $W$  to calculate accurate estimates with  $V$
- Another set of parameters  $\theta$  to estimate  $\pi$  given the output of  $V$ .
- Update both using gradient update with different learning rates

# Actor Critic

## Actor Critic Algorithm

- **Inputs:** differentiable  $\pi(a|s, \theta)$ , differentiable  $V(s, w)$
- **Parameters:** learning rates  $\eta_w, \eta_\theta > 0$ ,  $0 \leq \gamma \leq 1$
- **Variables:**  $\theta, w$  e.g. uniform initialization

**For** episode **in** `total_episodes`:

$t \leftarrow 0$ ,  $I \leftarrow 1$

$s_t \leftarrow \text{environment.reset}()$  # reset and get initial state

**While**  $t < T$  **and**  $s_t \neq s_{\text{terminal}}$ :

$a_t \leftarrow \pi(a|s, \theta)$

$s_{t+1}, r_t \leftarrow \text{environment.step}(a)$

$\delta \leftarrow r_t + \gamma V(s_{t+1}, w) - V(s_t, w)$

$w \leftarrow w + \eta_w \delta \nabla_w V(s_t, w)$

$\theta \leftarrow \theta + \eta_\theta \delta \nabla_\theta \ln[\pi(a|s_t, \theta)]$

$I \leftarrow \gamma I$ ,  $s_t \leftarrow s_{t+1}$

**Return**  $\pi(a|s, Q(s, a)), Q(s, a)$

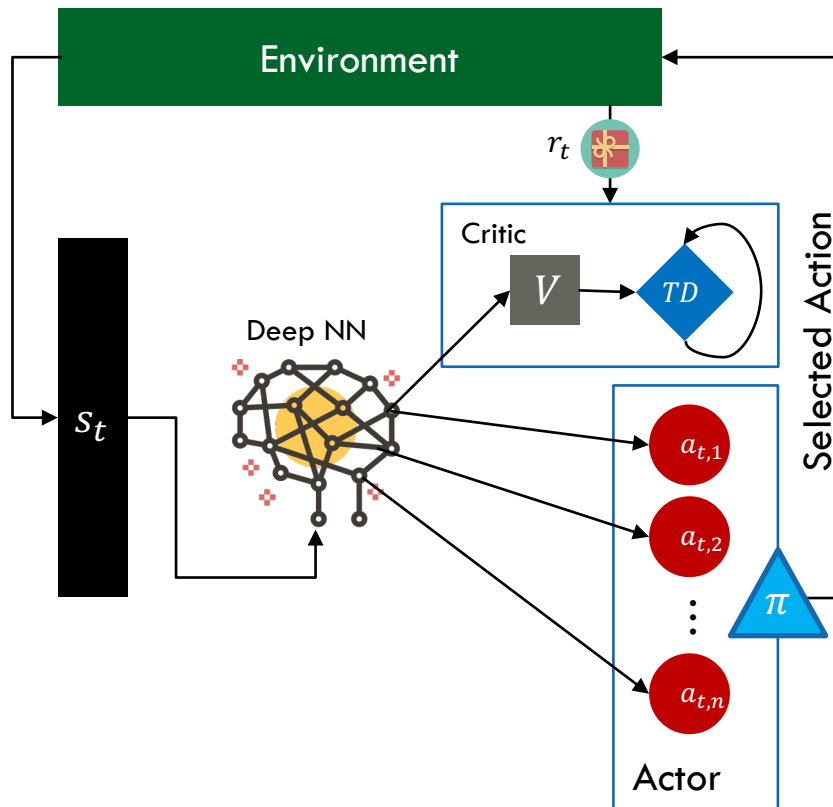
\* SUTTON, R. S. . B. (2018). *REINFORCEMENT LEARNING : an introduction*. MIT PRESS.  
Ch.13 Policy Gradient Methods, p.332

# Deep Learning

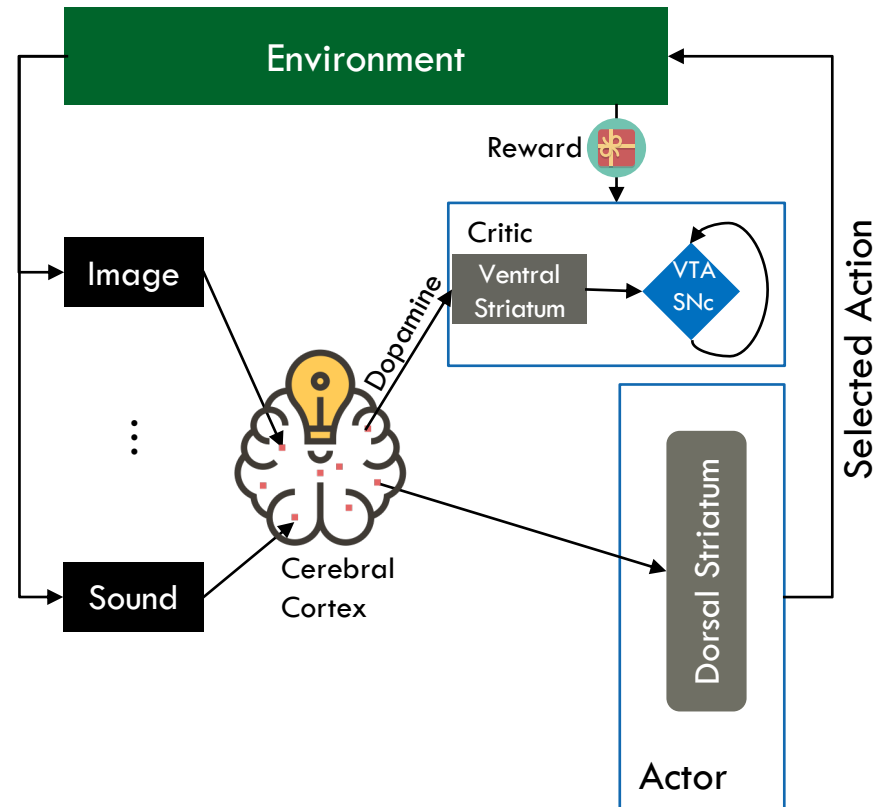
- Nice differentiable functions.
- High performance in learning tasks. Provides very generic and expressive approximators.
- Requires high data volumes  $\sim$  easily acquired in combinatorial problems, big data systems and real world applications.
- Efficient function approximation using non-linearity
- A variety of learning loss functions to be used as value  $\overline{VE}$  or policy  $J$  estimators

# Neural Actor-Critic

## Deep Learning Implementation



## Human Brain



# Multi-Agent Reinforcement Learning

## Independent

- A multi agent approach, where each agent learns to optimize their objectives independently of the others.
- An agent can interact with another agent by changing their environment and observing their actions via state changes.
- No extra interaction dynamics are modelled in this case.
- Usually it underperforms cooperative agents in cooperative tasks.

## Dependent:

- Shared Observations, e.g. a "super" agent that manages joint actions/observations.
- Aggregate over rewards
- Competitive rewards
- Latent Features Sharing

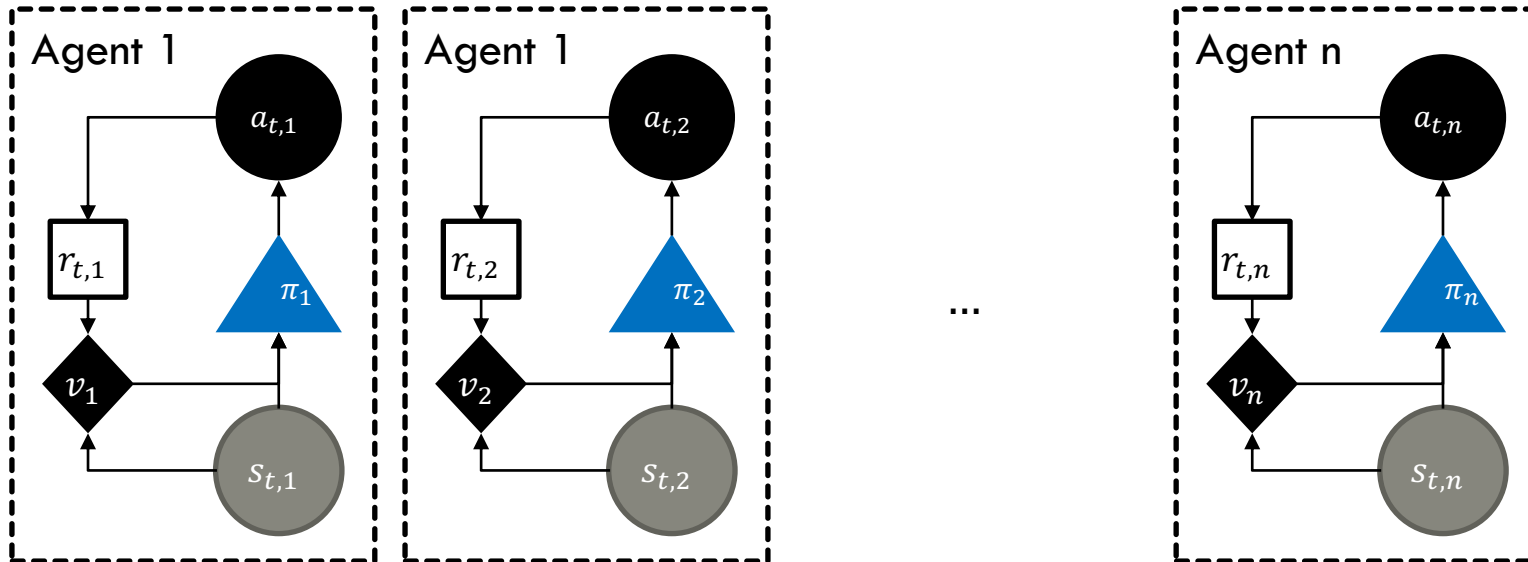
Tan, M. (1993). Multi-agent reinforcement learning: Independent vs. cooperative agents. In *Proceedings of the tenth international conference on machine learning* (pp. 330-337).



# Challenges of Multi-Agent Systems

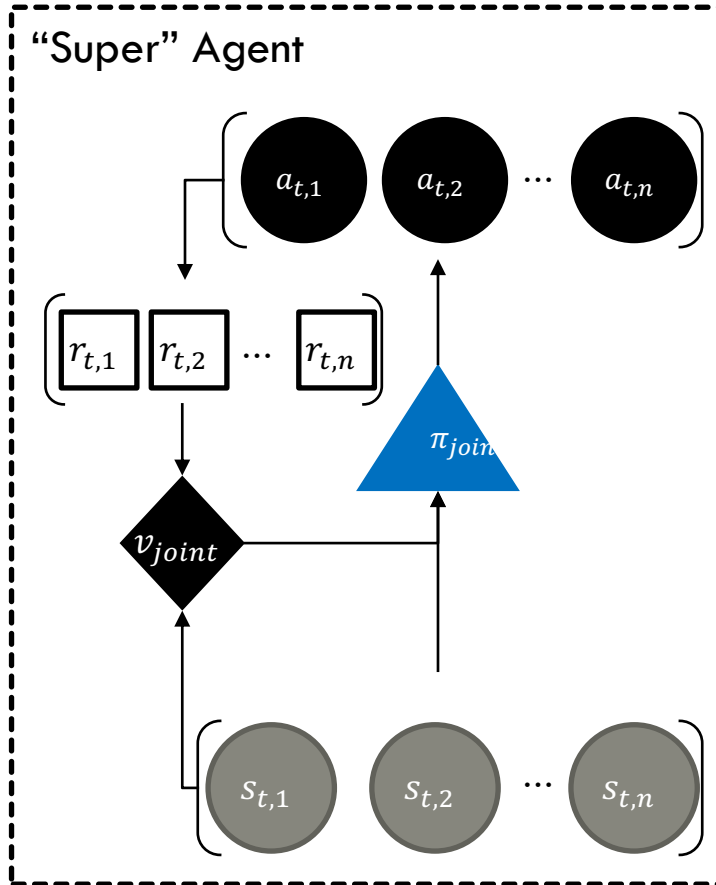
- Challenging Reward Design, e.g. average vs maxmin over all rewards?
- Individual Local observations → POMDP
- Heterogeneous agents → Difficult to model shared observations/rewards
- Scalability
- Small divergence in initial setting of agents usually diverges to much different states! (chaotic)

## Independent RL



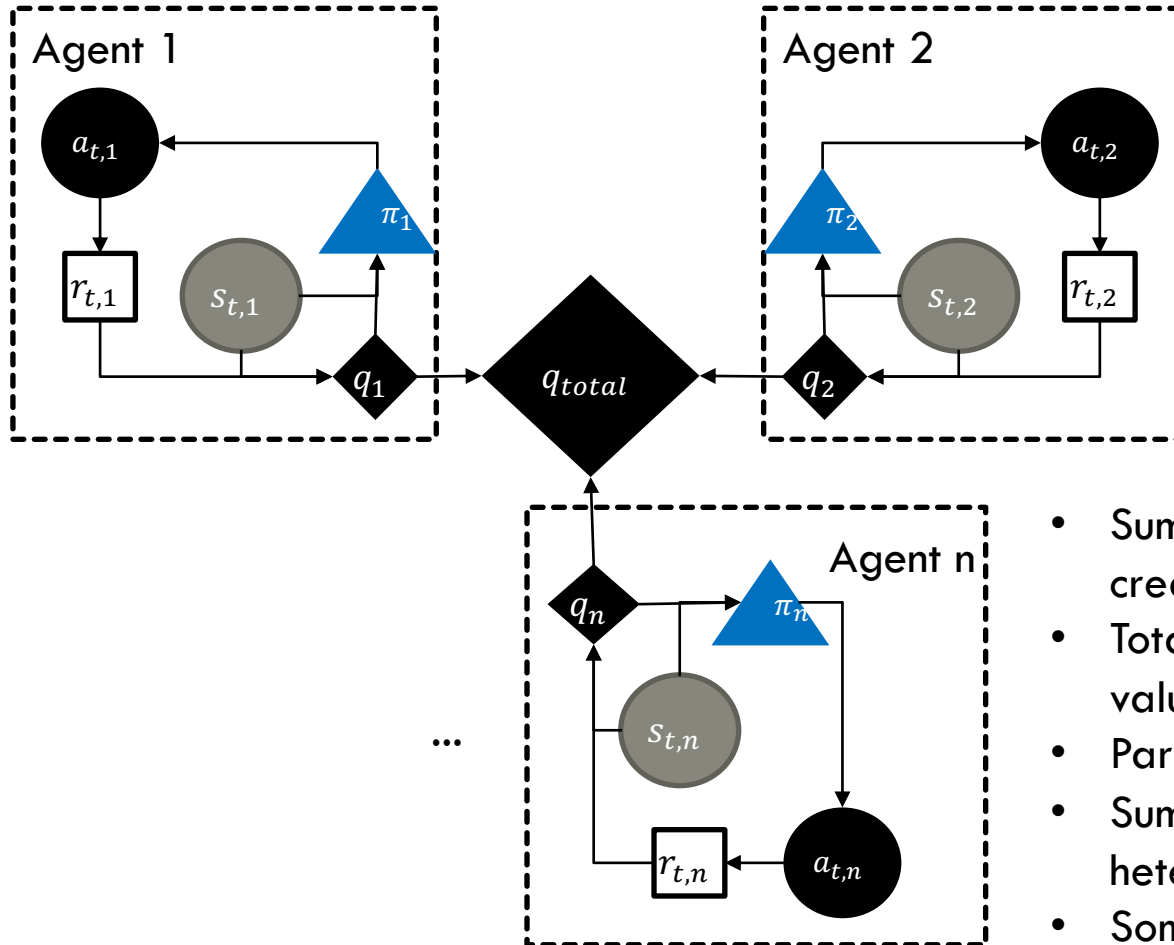
- Easiest Approach
- Expect that an agent infers behaviors of others by observation
- Scalable
- Works for heterogeneous agents
- Model Free (no communication modelling)
- In reality: doesn't learn agent interactions

# "Super" Agent



- Agent interactions learning within network
- Not scalable
- Centralized
- May be difficult to work for heterogeneous agents
- In reality: Rarely scales over 100 agents

# Value Decomposition Networks




- Summing individual values to create an aggregate one.
- Total q-value is a utility not a value
- Partly centralized
- Sum may not work for heterogeneous rewards
- Some Recent Architectures extend it: e.g. Q-Mix

# Summary


- Reinforcement Learning can be used for uncertain and dynamic environments
- Exploitation vs Exploration
- Control Problem: optimal policy,
- Prediction Problem: optimal estimation
- Deep learning can be used to tackle the prediction problem and state transformation for POMDP.
- Actor-Critic modelling tackles both prediction and control problems
- A simple approach to MARL can be done via Independent Reinforcement Learning.

# Questions



# Some References

- Sutton, R. S., & Barto, A. G. (1998). *Reinforcement learning : an introduction*. MIT Press. Retrieved from [https://drive.google.com/file/d/1opPSz5AZ\\_kVa1uWOdOiveNiBFiEOHjkG/view](https://drive.google.com/file/d/1opPSz5AZ_kVa1uWOdOiveNiBFiEOHjkG/view)  
 Interesting Chapters:
  - 1 Introduction: 1.1 Reinforcement Learning, 1.2 Examples, 1.3 Elements of Reinforcement Learning
  - 3 Finite Markov Processes: 3.1 The Agent-Environment Interface, 3.2 Goals and Rewards, 3.3 Returns and Episodes
  - 4 Dynamic Programming: All
  - 5 Monte Carlo Methods: All
  - 6 Temporal-Difference Learning: 6.1 TD Prediction, 6.2 Advantages of TD Prediction Methods, 6.4 Sarsa: On-policy TD Control, 6.5 Q-learning Off-policy TD Control
  - 8 Planning and Learning with Tabular Methods: 8.1 Models and Planning, 8.2 Dyna: Integrated Planning, Acting, and Learning, 8.13 Summary of Part I: Dimensions
  - 9 On-policy Prediction with Approximation: 9.1 Value-function Approximation, 9.2 The Prediction Objective ( $\overline{VE}$ ), 9.3 Stochastic-gradient and Semi-gradient Methods
  - 13 Policy Gradient Methods: 13.1 Policy Approximation and its Advantages, 13.2 The Policy Gradient Theorem, 13.5 Actor-Critic Methods
  - 14 Psychology: 14.1 Prediction and Control
  - 15 Neuroscience: 15.4 Dopamine, 15.7 Neural Actor-Critic, 15.9 Hedonistic Neurons, 15.10 Collective Reinforcement Learning
  - 17 Frontiers: 17.3 Observations and State
- Silver, D., Schrittwieser, J., Simonyan, K., Antonoglou, I., Huang, A., Guez, A., ... Hassabis, D. (2017). Mastering the game of Go without human knowledge. *Nature*, 550(7676), 354–359. <https://doi.org/10.1038/nature24270>
- David Silver's slides:  
<http://www0.cs.ucl.ac.uk/staff/d.silver/web/Teaching.html>



## Extra Slides

Just in case



# SARSA: On Policy Temporal Difference

$$Q(s_t, a_t) \leftarrow (1 - \alpha) Q(s_t, a_t) + \alpha \cdot (r_t + \gamma \cdot Q(s_{t+1}, a))$$

Learning rate

Discount factor

Reward

Value from previous calculation

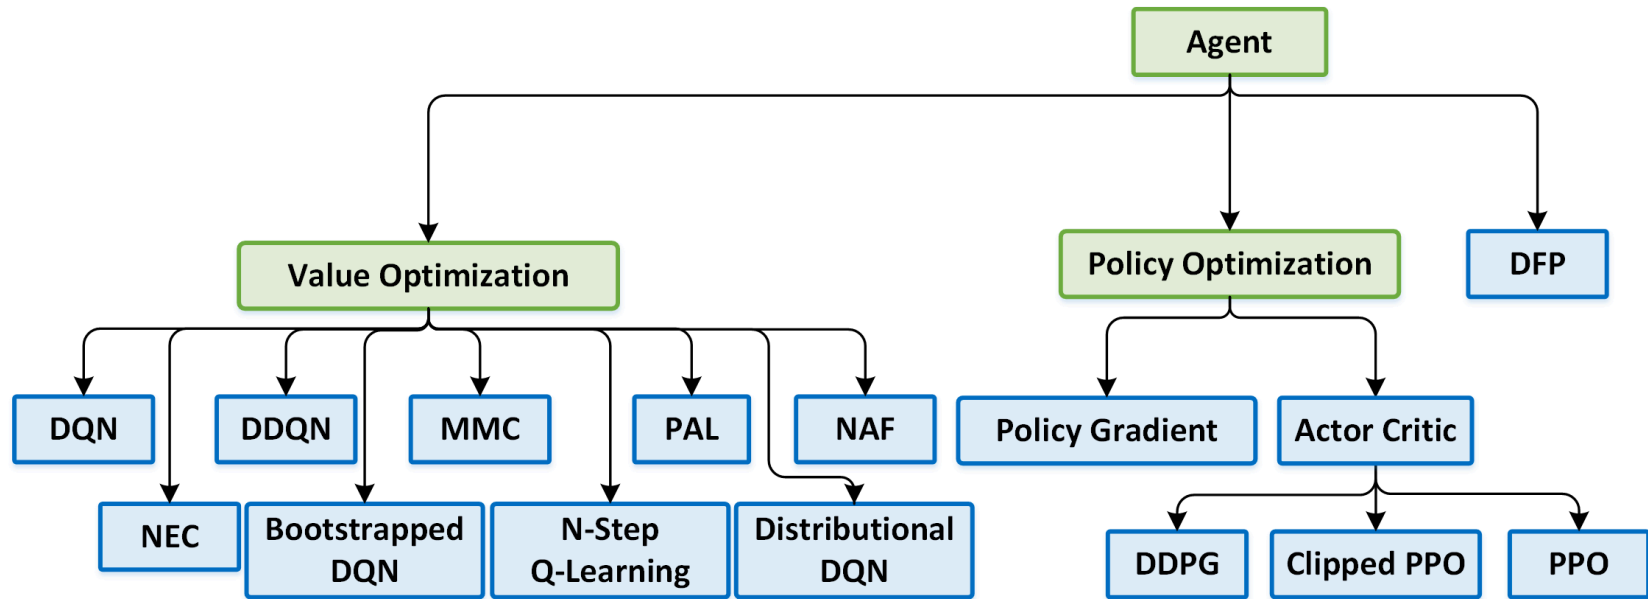
Estimate of optimal future value

The diagram illustrates the SARSA update equation. The equation is  $Q(s_t, a_t) \leftarrow (1 - \alpha) Q(s_t, a_t) + \alpha \cdot (r_t + \gamma \cdot Q(s_{t+1}, a))$ . Annotations include: a red arrow pointing to  $\alpha$  labeled 'Learning rate'; a green arrow pointing to  $r_t$  labeled 'Reward'; a red arrow pointing to  $\gamma$  labeled 'Discount factor'; a green arrow pointing down from  $Q(s_t, a_t)$  on the left side labeled 'Value from previous calculation'; and a green arrow pointing down from  $Q(s_{t+1}, a)$  labeled 'Estimate of optimal future value'.

# Training via Experience Buffer

- Allow a policy to act for some episodes or steps and then save each quadruplet  $(s_t, a_t, r_t, s_{t+1})$  in a buffer.
- Once enough buffers/samples are generated, make a training pass, by using  $s_t$  as input.
  - Select the target policy outputs  $\pi(a|\theta)$  corresponding to  $a_t$ .
  - Calculate  $V_\pi$  and  $Q_\pi$  based on  $r_t, s_{t+1}$  for all the steps in the buffer following  $s_t$ .
  - Calculate the gradients of value and policy estimators:  $\nabla_\theta J(\theta), \nabla_w \overline{VE}(w)$
  - Update parameters  $\theta, w$

# An RL Taxonomy



<https://ai.intel.com/reinforcement-learning-coach-intel/>

## Learning more

Some useful terms to check for learning more:

- Dueling networks
- Prioritized experience replay
- Advantage
- A2C
- Q-MIX
- Differentiable Inter-Agent Learning and Reinforced Inter-Agent Learning.

# Notation Table I

Symbol	Explanation
$i, j$	Agent indices
$O(x)$	An objective function that operates on input $x$
$t$	A timestep
$a_{t,i}$	An action taken by agent $i$ at time $t$
$s_{t,i}$	The agent state of agent $i$ at time $t$
$r_{t,i}$	The reward received by an agent $i$ at time $t$
$g(s_{t,i}a_{t,i})$	The state transition that happens from time $t$ to $t + 1$ given agent $i$ state and selected action
$V(s_{t,i})$	The value function, that provides the agent $i$ estimates about how optimal is its state at time $t$
$Q(s_{t,i}, a_{t,i})$	The action-value function, that provides the agent estimates about how optimal is its state $s_{t,i}$ and the action it selected $a_{t,i}$ at time $t$

## Notation Table II

Symbol	Explanation
$v(s_{t,i})$ $q(s_{t,i}, a_{t,i})$	The true state and action-state value functions that provided the actual value of how optimal the state $s_{t,i}$ and selected action $a_{t,i}$ are for agent $i$ at time $t$ . Usually, they are not known.
$\pi(a_{t,i} s_{t,i})$	The policy that selects the action $a_{t,i}$ given the state $s_{t,i}$ for the agent $i$ at time $t$ .
$\gamma$	The discount factor, usually $0 \leq \gamma \leq 1$ , which discounts future rewards and values
$R_{t,i}$	The cumulative reward from time $t$ and on, for agent $i$
$G_{t,i}$	The return, which is the cumulative reward from time $t$ until the end of an episode (e.g.. when a goal is met or failed for an agent $i$ ).
$\pi_*(a_{t,i} s_{t,i})$	The optimal policy that maximizes cumulative reward and return
$o_{t,i}$	The environmental observation of an agent $i$ at time $t$ . Usually modelled along with state.

## Notation Table III

Symbol	Explanation
$\alpha$	The learning rate in temporal difference models, usually $0 \leq \alpha \leq 1$
$\max_x f(x)$	The maximization of a function $f(x)$ in regards to $x$
$\min_x f(x)$	The minimization of a function $f(x)$ in regards to $x$
$w$	A tensor of learnable parameters for value estimation
$\overline{VE(w)}$	Value estimator of via parameters $w$
$\eta$	The learning rate for value and policy estimation via parameter learning
$\nabla_w f(w)$	The gradient of function $f(w)$ in regards to elements of $w$
$\mu(s)$	On-policy distribution. The amount of timesteps spent (or expected to be spent) on the state $s$
$\theta$	A tensor of learnable parameters for policy estimation
$J(\theta)$	Policy estimator via parameter learning